

Linear Classifiers (Part 3)

CS114B Lab 4

Kenneth Lai

February 17, 2022

Multi-Class Classification

- ▶ Two-class (binary) classification
 - ▶ Compute “score” $z = \theta \cdot \mathbf{x}$ (or $\mathbf{w} \cdot \mathbf{x} + b$)
 - ▶ Compute decision \hat{y} as a function of z

Multi-Class Classification

- ▶ Two-class (binary) classification
 - ▶ Compute “score” $z = \theta \cdot \mathbf{x}$ (or $\mathbf{w} \cdot \mathbf{x} + b$)
 - ▶ Compute decision \hat{y} as a function of z
 - ▶ If \hat{y} is interpreted as the probability of (or indicator for) one class, $1 - \hat{y}$ is the probability of (indicator for) the other class

Multi-Class Classification

- ▶ Two-class (binary) classification
 - ▶ Compute “score” $z = \theta \cdot \mathbf{x}$ (or $\mathbf{w} \cdot \mathbf{x} + b$)
 - ▶ Compute decision \hat{y} as a function of z
 - ▶ If \hat{y} is interpreted as the probability of (or indicator for) one class, $1 - \hat{y}$ is the probability of (indicator for) the other class
- ▶ Multi-class (multinomial) classification
 - ▶ Compute a **vector** of scores $\mathbf{z} = \Theta \cdot \mathbf{x}$ (or $\mathbf{W} \cdot \mathbf{x} + \mathbf{b}$)
 - ▶ Compute decision \hat{y} as a function of \mathbf{z}

Linear Maps and Matrices

- ▶ Given a feature vector \mathbf{x} , we want a score vector \mathbf{z}

Linear Maps and Matrices

- ▶ Given a feature vector \mathbf{x} , we want a score vector \mathbf{z}
- ▶ More generally, we want to define a **linear map** between the feature vector space and the score vector space

Linear Maps and Matrices

- ▶ Given a feature vector \mathbf{x} , we want a score vector \mathbf{z}
- ▶ More generally, we want to define a **linear map** between the feature vector space and the score vector space
- ▶ A **matrix** is a rectangular array of scalars, that can be used to define a linear map

Linear Maps and Matrices

- ▶ Warning! A note on notation:

Linear Maps and Matrices

- ▶ Warning! A note on notation:
 - ▶ Math convention: A p -by- n matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p

Linear Maps and Matrices

- ▶ Warning! A note on notation:
 - ▶ Math convention: A p -by- n matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p
 - ▶ Matrices have shapes (output dimension, input dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{p \times n}$, $\mathbf{x} \in \mathbb{R}^n$, and $\mathbf{z} \in \mathbb{R}^p$
 - ▶ $\mathbf{z} = \Theta \cdot \mathbf{x}$ (or $\mathbf{W} \cdot \mathbf{x} + \mathbf{b}$)

Linear Maps and Matrices

- ▶ Warning! A note on notation:
 - ▶ Math convention: A p -by- n matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p
 - ▶ Matrices have shapes (output dimension, input dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{p \times n}$, $\mathbf{x} \in \mathbb{R}^n$, and $\mathbf{z} \in \mathbb{R}^p$
 - ▶ $\mathbf{z} = \Theta \cdot \mathbf{x}$ (or $\mathbf{W} \cdot \mathbf{x} + \mathbf{b}$)
 - ▶ Computer science convention (mostly): An n -by- p matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p (technically $\mathbb{R}^{1 \times n}$ to $\mathbb{R}^{1 \times p}$)

Linear Maps and Matrices

- ▶ Warning! A note on notation:
 - ▶ Math convention: A p -by- n matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p
 - ▶ Matrices have shapes (output dimension, input dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{p \times n}$, $\mathbf{x} \in \mathbb{R}^n$, and $\mathbf{z} \in \mathbb{R}^p$
 - ▶ $\mathbf{z} = \Theta \cdot \mathbf{x}$ (or $\mathbf{W} \cdot \mathbf{x} + \mathbf{b}$)
 - ▶ Computer science convention (mostly): An n -by- p matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p (technically $\mathbb{R}^{1 \times n}$ to $\mathbb{R}^{1 \times p}$)
 - ▶ Matrices have shapes (input dimension, output dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{n \times p}$, $\mathbf{x} \in \mathbb{R}^{1 \times n}$, and $\mathbf{z} \in \mathbb{R}^{1 \times p}$
 - ▶ $\mathbf{z} = \mathbf{x} \cdot \Theta$ (or $\mathbf{x} \cdot \mathbf{W} + \mathbf{b}$)

Linear Maps and Matrices

- ▶ Warning! A note on notation:
 - ▶ Math convention: A p -by- n matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p
 - ▶ Matrices have shapes (output dimension, input dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{p \times n}$, $\mathbf{x} \in \mathbb{R}^n$, and $\mathbf{z} \in \mathbb{R}^p$
 - ▶ $\mathbf{z} = \Theta \cdot \mathbf{x}$ (or $\mathbf{W} \cdot \mathbf{x} + \mathbf{b}$)
 - ▶ Computer science convention (mostly): An n -by- p matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p (technically $\mathbb{R}^{1 \times n}$ to $\mathbb{R}^{1 \times p}$)
 - ▶ Matrices have shapes (input dimension, output dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{n \times p}$, $\mathbf{x} \in \mathbb{R}^{1 \times n}$, and $\mathbf{z} \in \mathbb{R}^{1 \times p}$
 - ▶ $\mathbf{z} = \mathbf{x} \cdot \Theta$ (or $\mathbf{x} \cdot \mathbf{W} + \mathbf{b}$)
 - ▶ More intuitive (input \rightarrow output)

Linear Maps and Matrices

- ▶ Warning! A note on notation:
 - ▶ Math convention: A p -by- n matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p
 - ▶ Matrices have shapes (output dimension, input dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{p \times n}$, $\mathbf{x} \in \mathbb{R}^n$, and $\mathbf{z} \in \mathbb{R}^p$
 - ▶ $\mathbf{z} = \Theta \cdot \mathbf{x}$ (or $\mathbf{W} \cdot \mathbf{x} + \mathbf{b}$)
 - ▶ Computer science convention (mostly): An n -by- p matrix defines a linear map from \mathbb{R}^n to \mathbb{R}^p (technically $\mathbb{R}^{1 \times n}$ to $\mathbb{R}^{1 \times p}$)
 - ▶ Matrices have shapes (input dimension, output dimension)
 - ▶ Let $\Theta \in \mathbb{R}^{n \times p}$, $\mathbf{x} \in \mathbb{R}^{1 \times n}$, and $\mathbf{z} \in \mathbb{R}^{1 \times p}$
 - ▶ $\mathbf{z} = \mathbf{x} \cdot \Theta$ (or $\mathbf{x} \cdot \mathbf{W} + \mathbf{b}$)
 - ▶ More intuitive (input \rightarrow output)
 - ▶ Aligns with the convention in (mini)batch training that the first dimension is the batch size ("feature vectors are stacked row-wise")

General Advice

- ▶ Know your shapes!

General Advice

- ▶ Know your shapes!



Source

Multi-Class Classification

- ▶ What is the activation function g ?

Multi-Class Classification

- ▶ What is the activation function g ?
- ▶ Let $\mathbf{z} = [z_1 \ \dots \ z_p]$ be a vector of scores for each class

Multi-Class Classification

- ▶ What is the activation function g ?
- ▶ Let $\mathbf{z} = [z_1 \ \dots \ z_p]$ be a vector of scores for each class
- ▶ Logistic regression: softmax function
 - ▶ $\text{softmax}(z_c) = \frac{e^{z_c}}{\sum_{k=1}^p e^{z_k}} = P(y = c|\mathbf{x}) = \hat{y}_c$

Multi-Class Classification

- ▶ What is the activation function g ?
- ▶ Let $\mathbf{z} = [z_1 \ \dots \ z_p]$ be a vector of scores for each class
- ▶ Logistic regression: softmax function
 - ▶ $\text{softmax}(z_c) = \frac{e^{z_c}}{\sum_{k=1}^p e^{z_k}} = P(y = c|\mathbf{x}) = \hat{y}_c$
 - ▶ $\text{softmax}(\mathbf{z}) = \hat{\mathbf{y}}$ is a vector of probabilities for each class

Multi-Class Classification

- ▶ What is the activation function g ?
- ▶ Let $\mathbf{z} = [z_1 \ \dots \ z_p]$ be a vector of scores for each class
- ▶ Logistic regression: softmax function
 - ▶ $\text{softmax}(z_c) = \frac{e^{z_c}}{\sum_{k=1}^p e^{z_k}} = P(y = c|\mathbf{x}) = \hat{y}_c$
 - ▶ $\text{softmax}(\mathbf{z}) = \hat{\mathbf{y}}$ is a vector of probabilities for each class
- ▶ Perceptron: argmax function
 - ▶ $\underset{k=1}{\text{argmax}}^p(z_k) = \hat{y}$

Multi-Class Classification

- ▶ What is the activation function g ?
- ▶ Let $\mathbf{z} = [z_1 \ \dots \ z_p]$ be a vector of scores for each class
- ▶ Logistic regression: softmax function
 - ▶ $\text{softmax}(z_c) = \frac{e^{z_c}}{\sum_{k=1}^p e^{z_k}} = P(y = c|\mathbf{x}) = \hat{y}_c$
 - ▶ $\text{softmax}(\mathbf{z}) = \hat{\mathbf{y}}$ is a vector of probabilities for each class
- ▶ Perceptron: argmax function
 - ▶ $\text{argmax}_{k=1}^p(z_k) = \hat{y}$
 - ▶ What if there is a tie?
 - ▶ Do whatever `numpy.argmax` does

Naïve Bayes as a Linear Classifier

- ▶ Suppose we observe a document d . What is the most likely class \hat{c} ?

Naïve Bayes as a Linear Classifier

- ▶ Suppose we observe a document d . What is the most likely class \hat{c} ?

- ▶
$$P(c|d) = \frac{P(d|c)P(c)}{P(d)}$$

- ▶ Bayes' Rule

- ▶
$$\hat{c} = \operatorname{argmax}_{c \in C} P(d|c)P(c)$$

- ▶ $P(d)$ is the same for each class

- ▶
$$\hat{c} = \operatorname{argmax}_{c \in C} P(c) \prod_{i \in \text{positions}} P(w_i|c)$$

- ▶ Bag of words assumption, Naïve Bayes assumption

- ▶
$$\hat{c} = \operatorname{argmax}_{c \in C} \log P(c) + \sum_{i \in \text{positions}} \log P(w_i|c)$$

- ▶ If $xy = z$, then $\log(x) + \log(y) = \log(z)$

Naïve Bayes as a Linear Classifier

- ▶ $\hat{c} = \operatorname{argmax}_{c \in C} \sum_{w \in |V|} \left[(\log P(w|c))(\operatorname{count}(w, d)) \right] + \log P(c)$
- ▶ $\sum_{i \in \text{positions}} \log P(w_i|c) = \sum_{w \in |V|} \left[(\log P(w|c))(\operatorname{count}(w, d)) \right]$

Naïve Bayes as a Linear Classifier

- ▶ $\hat{c} = \operatorname{argmax}_{c \in C} \sum_{w \in |V|} \left[(\log P(w|c))(\operatorname{count}(w, d)) \right] + \log P(c)$
 - ▶ $\sum_{i \in \text{positions}} \log P(w_i|c) = \sum_{w \in |V|} \left[(\log P(w|c))(\operatorname{count}(w, d)) \right]$
- ▶ Let $\ell_{cw} = \log P(w|c)$, $x_w = \operatorname{count}(w, d)$, and $p_c = \log P(c)$

Naïve Bayes as a Linear Classifier

- ▶ $\hat{c} = \operatorname{argmax}_{c \in C} \sum_{w \in |V|} \left[(\log P(w|c))(\operatorname{count}(w, d)) \right] + \log P(c)$
 - ▶ $\sum_{i \in \text{positions}} \log P(w_i|c) = \sum_{w \in |V|} \left[(\log P(w|c))(\operatorname{count}(w, d)) \right]$
- ▶ Let $\ell_{cw} = \log P(w|c)$, $x_w = \operatorname{count}(w, d)$, and $p_c = \log P(c)$

$$\hat{c} = \operatorname{argmax}_{c \in C} \sum_{w \in |V|} \ell_{cw} x_w + p_c$$

Naïve Bayes as a Linear Classifier

- ▶ $\hat{c} = \operatorname{argmax}_{c \in C} \sum_{w \in |V|} \left[(\log P(w|c))(\operatorname{count}(w, d)) \right] + \log P(c)$
 - ▶ $\sum_{i \in \text{positions}} \log P(w_i|c) = \sum_{w \in |V|} \left[(\log P(w|c))(\operatorname{count}(w, d)) \right]$
- ▶ Let $\ell_{cw} = \log P(w|c)$, $x_w = \operatorname{count}(w, d)$, and $p_c = \log P(c)$

$$\begin{aligned}\hat{c} &= \operatorname{argmax}_{c \in C} \sum_{w \in |V|} \ell_{cw} x_w + p_c \\ &= \operatorname{argmax}_{c \in C} (\ell_c \cdot \mathbf{x} + p_c)\end{aligned}$$

Naïve Bayes as a Linear Classifier

- ▶ Let \mathbf{x} be a feature vector, $\mathbf{p} = \text{self.prior}$, and $\mathcal{L} = \text{self.likelihood}$

Naïve Bayes as a Linear Classifier

- ▶ Let \mathbf{x} be a feature vector, $\mathbf{p} = \text{self.prior}$, and $\mathcal{L} = \text{self.likelihood}$
- ▶ $\mathbf{z} = \mathcal{L} \cdot \mathbf{x} + \mathbf{p}$

Naïve Bayes as a Linear Classifier

- ▶ Let \mathbf{x} be a feature vector, $\mathbf{p} = \text{self.prior}$, and $\mathcal{L} = \text{self.likelihood}$
- ▶ $\mathbf{z} = \mathcal{L} \cdot \mathbf{x} + \mathbf{p}$
- ▶ $\hat{c} = \underset{c \in C}{\operatorname{argmax}}(z_c)$

Training Linear Classifiers

- ▶ Naïve Bayes: estimate parameters (log-prior, log-likelihood) directly from training data

Training Linear Classifiers

- ▶ Naïve Bayes: estimate parameters (log-prior, log-likelihood) directly from training data
- ▶ Logistic regression, perceptron:
 - ▶ Define a **loss function**

Training Linear Classifiers

- ▶ Naïve Bayes: estimate parameters (log-prior, log-likelihood) directly from training data
- ▶ Logistic regression, perceptron:
 - ▶ Define a **loss function**
 - ▶ Update the parameters using **gradient descent**

Loss Functions

- ▶ How wrong is your classifier?

Loss Functions

- ▶ How wrong is your classifier?
- ▶ Logistic regression: cross-entropy loss
 - ▶ $L(\hat{y}, y) = -\log P(y|\mathbf{x}) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$

Loss Functions

- ▶ How wrong is your classifier?
- ▶ Logistic regression: cross-entropy loss
 - ▶ $L(\hat{y}, y) = -\log P(y|\mathbf{x}) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$
 - ▶ Minimizing loss = maximizing the (log-)probability of the true y given \mathbf{x}

Loss Functions

- ▶ How wrong is your classifier?
- ▶ Logistic regression: cross-entropy loss
 - ▶ $L(\hat{y}, y) = -\log P(y|\mathbf{x}) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$
 - ▶ Minimizing loss = maximizing the (log-)probability of the true y given \mathbf{x}
- ▶ Perceptron: perceptron loss
 - ▶ $L(\hat{y}, y) = (\hat{y} - y)z$ (for $y \in \{0, 1\}$)

Loss Functions

- ▶ How wrong is your classifier?
- ▶ Logistic regression: cross-entropy loss
 - ▶ $L(\hat{y}, y) = -\log P(y|\mathbf{x}) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$
 - ▶ Minimizing loss = maximizing the (log-)probability of the true y given \mathbf{x}
- ▶ Perceptron: perceptron loss
 - ▶ $L(\hat{y}, y) = (\hat{y} - y)z$ (for $y \in \{0, 1\}$)
 - ▶ (You may also see $L = \max(0, -yz)$, for $y \in \{-1, 1\}$)

Loss Functions

- ▶ How wrong is your classifier?
- ▶ Logistic regression: cross-entropy loss
 - ▶ $L(\hat{y}, y) = -\log P(y|\mathbf{x}) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$
 - ▶ Minimizing loss = maximizing the (log-)probability of the true y given \mathbf{x}
- ▶ Perceptron: perceptron loss
 - ▶ $L(\hat{y}, y) = (\hat{y} - y)z$ (for $y \in \{0, 1\}$)
 - ▶ (You may also see $L = \max(0, -yz)$, for $y \in \{-1, 1\}$)
 - ▶ If $\hat{y} = y$, $L = 0$
 - ▶ If $\hat{y} \neq y$, $L > 0$

Derivatives

- ▶ The **derivative** of a function measures the instantaneous rate of change in a function's output with respect to a change in its input

Derivatives

- ▶ The **derivative** of a function measures the instantaneous rate of change in a function's output with respect to a change in its input
 - ▶ "Slope of a function's graph"

Derivatives

- ▶ The **derivative** of a function measures the instantaneous rate of change in a function's output with respect to a change in its input
 - ▶ “Slope of a function's graph”
- ▶ The derivative of f with respect to x is denoted $\frac{df}{dx}$, f' , etc.

Derivatives

► Rules of differentiation

- Constant rule: If $f(x)$ is constant, then $f'(x) = 0$
- Sum rule: $(f + g)' = f' + g'$
- Product rule: $(fg)' = f'g + fg'$
- Power rule: If $f(x) = x^r$, then $f'(x) = rx^{r-1}$
- Chain rule: If $h(x) = f(g(x))$, then $h'(x) = f'(g(x)) \cdot g'(x)$
(or $\frac{dh}{dx} = \frac{df}{dg} \cdot \frac{dg}{dx}$)

Derivatives

- ▶ Rules of differentiation
 - ▶ Constant rule: If $f(x)$ is constant, then $f'(x) = 0$
 - ▶ Sum rule: $(f + g)' = f' + g'$
 - ▶ Product rule: $(fg)' = f'g + fg'$
 - ▶ Power rule: If $f(x) = x^r$, then $f'(x) = rx^{r-1}$
 - ▶ Chain rule: If $h(x) = f(g(x))$, then $h'(x) = f'(g(x)) \cdot g'(x)$
(or $\frac{dh}{dx} = \frac{df}{dg} \cdot \frac{dg}{dx}$)
- ▶ Anything more complicated than this, we will tell you what the derivative is

Partial Derivatives

- ▶ The **partial derivative** of a function **of several variables** measures the instantaneous rate of change in a function's output with respect to a change in **one of** its inputs, **with the others held constant**

Partial Derivatives

- ▶ The **partial derivative** of a function **of several variables** measures the instantaneous rate of change in a function's output with respect to a change in **one of** its inputs, **with the others held constant**
- ▶ The partial derivative of f with respect to x is denoted $\frac{\partial f}{\partial x}$, f_x , etc.

Gradients

- ▶ The **gradient** of a function of several variables is a vector of partial derivatives

Gradients

- ▶ The **gradient** of a function of several variables is a vector of partial derivatives

$$\nabla F = \begin{bmatrix} \frac{\partial F}{\partial x_1} \\ \vdots \\ \frac{\partial F}{\partial x_n} \end{bmatrix}$$

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L

$$\nabla L = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \vdots \\ \frac{\partial L}{\partial w_n} \\ \frac{\partial L}{\partial b} \end{bmatrix}$$

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L

$$\nabla L = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \vdots \\ \frac{\partial L}{\partial w_n} \\ \frac{\partial L}{\partial b} \end{bmatrix}$$

- ▶ \approx slope of loss function

Gradient Descent




Figure 5.4 The first step in iteratively finding the minimum of this loss function, by moving w in the reverse direction from the slope of the function. Since the slope is negative, we need to move w in a positive direction, to the right. Here superscripts are used for learning steps, so w^1 means the initial value of w (which is 0), w^2 at the second step, and so on.

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$
 - ▶ η = learning rate

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$
 - ▶ η = learning rate
 - ▶ “Hyperparameter”: parameter set before training

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$
 - ▶ η = learning rate
 - ▶ “Hyperparameter”: parameter set before training
 - ▶ Trade-off between speed of convergence and “zig-zag” behavior

Gradient Descent

- ▶ Initialize parameters $\theta = \mathbf{w}, b$ (randomly or $\mathbf{0}$)
- ▶ At each time step t :
 - ▶ Compute gradient ∇L
 - ▶ Move in direction of negative gradient
- ▶ $\theta_{t+1} = \theta_t - \eta \nabla L$
 - ▶ η = learning rate
 - ▶ “Hyperparameter”: parameter set before training
 - ▶ Trade-off between speed of convergence and “zig-zag” behavior
 - ▶ Often a function of t

Gradients in Logistic Regression

► $L(\hat{y}, y) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$

Gradients in Logistic Regression

- ▶ $L(\hat{y}, y) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$
- ▶ ...
- ▶ (calculus—see supplement slides)
- ▶ ...

Gradients in Logistic Regression

- ▶ $L(\hat{y}, y) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$
- ▶ ...
- ▶ (calculus—see supplement slides)
- ▶ ...
- ▶ $\frac{\partial L}{\partial w_j} = (\hat{y} - y)x_j$
- ▶ $\frac{\partial L}{\partial b} = \hat{y} - y$

Gradients in Logistic Regression

- ▶ $L(\hat{y}, y) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$
- ▶ ...
- ▶ (calculus—see supplement slides)
- ▶ ...
- ▶ $\frac{\partial L}{\partial w_j} = (\hat{y} - y)x_j$
- ▶ $\frac{\partial L}{\partial b} = \hat{y} - y$

- ▶ $\nabla L = (\hat{y} - y)\mathbf{x}$

Gradients in Perceptrons

► $L(\hat{y}, y) = (\hat{y} - y)z$

Gradients in Perceptrons

- ▶ $L(\hat{y}, y) = (\hat{y} - y)z$
- ▶ ...
- ▶ (calculus—see supplement slides)
- ▶ ...

Gradients in Perceptrons

- ▶ $L(\hat{y}, y) = (\hat{y} - y)z$
- ▶ ...
- ▶ (calculus—see supplement slides)
- ▶ ...
- ▶ $\frac{\partial L}{\partial w_j} = (\hat{y} - y)x_j$
- ▶ $\frac{\partial L}{\partial b} = \hat{y} - y$

Gradients in Perceptrons

- ▶ $L(\hat{y}, y) = (\hat{y} - y)z$
- ▶ ...
- ▶ (calculus—see supplement slides)
- ▶ ...
- ▶ $\frac{\partial L}{\partial w_j} = (\hat{y} - y)x_j$
- ▶ $\frac{\partial L}{\partial b} = \hat{y} - y$

- ▶ $\nabla L = (\hat{y} - y)\mathbf{x}$

Gradients in Perceptrons

- ▶ $L(\hat{y}, y) = (\hat{y} - y)z$
- ▶ ...
- ▶ (calculus—see supplement slides)
- ▶ ...
- ▶ $\frac{\partial L}{\partial w_j} = (\hat{y} - y)x_j$
- ▶ $\frac{\partial L}{\partial b} = \hat{y} - y$

- ▶ $\nabla L = (\hat{y} - y)\mathbf{x}$

- ▶ Does this look familiar?

Perceptron Learning Algorithm

- ▶ If $\hat{y} = y$, then $\hat{y} - y = 0$
 - ▶ Do nothing

Perceptron Learning Algorithm

- ▶ If $\hat{y} = y$, then $\hat{y} - y = 0$
 - ▶ Do nothing
- ▶ If $\hat{y} = 0$ and $y = 1$, then $\hat{y} - y = -1$
 - ▶ $\nabla L = -\mathbf{x}$
 - ▶ $\theta_{t+1} = \theta_t + \eta \mathbf{x}$
 - ▶ Increment weights

Perceptron Learning Algorithm

- ▶ If $\hat{y} = y$, then $\hat{y} - y = 0$
 - ▶ Do nothing
- ▶ If $\hat{y} = 0$ and $y = 1$, then $\hat{y} - y = -1$
 - ▶ $\nabla L = -\mathbf{x}$
 - ▶ $\theta_{t+1} = \theta_t + \eta \mathbf{x}$
 - ▶ Increment weights
- ▶ If $\hat{y} = 1$ and $y = 0$, then $\hat{y} - y = 1$
 - ▶ $\nabla L = \mathbf{x}$
 - ▶ $\theta_{t+1} = \theta_t - \eta \mathbf{x}$
 - ▶ Decrement weights

Gradients in Multinomial Logistic Regression

► Cross-entropy loss $L(\hat{\mathbf{y}}, \mathbf{y}) = - \sum_{k=1}^p y_k \log \hat{y}_k$

Gradients in Multinomial Logistic Regression

- ▶ Cross-entropy loss $L(\hat{\mathbf{y}}, \mathbf{y}) = - \sum_{k=1}^p y_k \log \hat{y}_k$
- ▶ Gradient ∇L becomes a matrix, where
 - ▶ $\frac{\partial L}{\partial w_{jk}} = (\hat{y}_k - y_k)x_j$
 - ▶ $\frac{\partial L}{\partial b_k} = \hat{y}_k - y_k$

Gradients in Multinomial Logistic Regression

- ▶ Cross-entropy loss $L(\hat{\mathbf{y}}, \mathbf{y}) = - \sum_{k=1}^p y_k \log \hat{y}_k$
- ▶ Gradient ∇L becomes a matrix, where
 - ▶ $\frac{\partial L}{\partial w_{jk}} = (\hat{y}_k - y_k)x_j$
 - ▶ $\frac{\partial L}{\partial b_k} = \hat{y}_k - y_k$
- ▶ $\nabla L = \mathbf{x} \otimes (\hat{\mathbf{y}} - \mathbf{y})$, where
 - ▶ \otimes denotes the outer product

Multi-class Perceptron Learning Algorithm

- ▶ Multi-class perceptron loss $L(\hat{y}, y) = z_{\hat{y}} - z_y$

Multi-class Perceptron Learning Algorithm

- ▶ Multi-class perceptron loss $L(\hat{y}, y) = z_{\hat{y}} - z_y$
- ▶ If $\hat{y} = y$, then do nothing
- ▶ Else:

Multi-class Perceptron Learning Algorithm

- ▶ Multi-class perceptron loss $L(\hat{y}, y) = z_{\hat{y}} - z_y$
- ▶ If $\hat{y} = y$, then do nothing
- ▶ Else:
 - ▶ For the correct class y , $\frac{\partial L}{\partial z_y} = -1$
 - ▶ $(\nabla L)_y = -\mathbf{x}$
 - ▶ $(\theta_y)_{t+1} = (\theta_y)_t + \eta \mathbf{x}$
 - ▶ Increment weights

Multi-class Perceptron Learning Algorithm

- ▶ Multi-class perceptron loss $L(\hat{y}, y) = z_{\hat{y}} - z_y$
- ▶ If $\hat{y} = y$, then do nothing
- ▶ Else:
 - ▶ For the correct class y , $\frac{\partial L}{\partial z_y} = -1$
 - ▶ $(\nabla L)_y = -\mathbf{x}$
 - ▶ $(\theta_y)_{t+1} = (\theta_y)_t + \eta \mathbf{x}$
 - ▶ Increment weights
 - ▶ For the predicted class \hat{y} , $\frac{\partial L}{\partial z_{\hat{y}}} = 1$
 - ▶ $(\nabla L)_{\hat{y}} = \mathbf{x}$
 - ▶ $(\theta_{\hat{y}})_{t+1} = (\theta_{\hat{y}})_t - \eta \mathbf{x}$
 - ▶ Decrement weights

Multi-class Perceptron Learning Algorithm

- ▶ Multi-class perceptron loss $L(\hat{y}, y) = z_{\hat{y}} - z_y$
- ▶ If $\hat{y} = y$, then do nothing
- ▶ Else:
 - ▶ For the correct class y , $\frac{\partial L}{\partial z_y} = -1$
 - ▶ $(\nabla L)_y = -\mathbf{x}$
 - ▶ $(\theta_y)_{t+1} = (\theta_y)_t + \eta \mathbf{x}$
 - ▶ Increment weights
 - ▶ For the predicted class \hat{y} , $\frac{\partial L}{\partial z_{\hat{y}}} = 1$
 - ▶ $(\nabla L)_{\hat{y}} = \mathbf{x}$
 - ▶ $(\theta_{\hat{y}})_{t+1} = (\theta_{\hat{y}})_t - \eta \mathbf{x}$
 - ▶ Decrement weights
 - ▶ For other classes, do nothing

(Mini-)Batch Training

- ▶ What is a time step?

(Mini-)Batch Training

- ▶ What is a time step?
 - ▶ Stochastic gradient descent: update θ after every training example

(Mini-)Batch Training

- ▶ What is a time step?
 - ▶ Stochastic gradient descent: update θ after every training example
 - ▶ Can result in very choppy movements

(Mini-)Batch Training

- ▶ What is a time step?
 - ▶ Stochastic gradient descent: update θ after every training example
 - ▶ Can result in very choppy movements
 - ▶ Batch gradient descent: update θ after processing the entire training set

(Mini-)Batch Training

- ▶ What is a time step?
 - ▶ Stochastic gradient descent: update θ after every training example
 - ▶ Can result in very choppy movements
 - ▶ Batch gradient descent: update θ after processing the entire training set
 - ▶ Minibatch gradient descent: update θ after m training examples

(Mini-)Batch Training

- ▶ What is a time step?
 - ▶ Stochastic gradient descent: update θ after every training example
 - ▶ Can result in very choppy movements
 - ▶ Batch gradient descent: update θ after processing the entire training set
 - ▶ Minibatch gradient descent: update θ after m training examples
 - ▶ Gradient = average of individual gradients

(Mini-)Batch Training

- ▶ Let \mathbf{x} consist of the feature vectors $\mathbf{x}^{(i)}$ for each document i in the (mini-)batch of size m , stacked on top of each other

(Mini-)Batch Training

- ▶ Let \mathbf{x} consist of the feature vectors $\mathbf{x}^{(i)}$ for each document i in the (mini-)batch of size m , stacked on top of each other

- ▶ $\mathbf{x} = \begin{bmatrix} x_1^{(1)} & \dots & x_n^{(1)} & 1 \\ \vdots & \ddots & \vdots & \vdots \\ x_1^{(m)} & \dots & x_n^{(m)} & 1 \end{bmatrix}$

(Mini-)Batch Training

- ▶ Let \mathbf{x} consist of the feature vectors $\mathbf{x}^{(i)}$ for each document i in the (mini-)batch of size m , stacked on top of each other

- ▶ $\mathbf{x} = \begin{bmatrix} x_1^{(1)} & \dots & x_n^{(1)} & 1 \\ \vdots & \ddots & \vdots & \vdots \\ x_1^{(m)} & \dots & x_n^{(m)} & 1 \end{bmatrix}$

- ▶ $\mathbf{y} = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(m)} \end{bmatrix}$

(Mini-)Batch Training

$$\blacktriangleright \nabla L = \frac{1}{m} (\mathbf{x}^T \cdot (\hat{\mathbf{y}} - \mathbf{y}))$$

(Mini-)Batch Training

- ▶ $\nabla L = \frac{1}{m}(\mathbf{x}^T \cdot (\hat{\mathbf{y}} - \mathbf{y}))$
 - ▶ What is $\mathbf{x}^T \cdot (\hat{\mathbf{y}} - \mathbf{y})$?

(Mini-)Batch Training

$$\begin{bmatrix} x_1^{(1)} & \dots & x_1^{(m)} \\ \vdots & \ddots & \vdots \\ x_n^{(1)} & \dots & x_n^{(m)} \\ 1 & 1 & 1 \end{bmatrix} \cdot \begin{bmatrix} \hat{y}^{(1)} - y^{(1)} \\ \vdots \\ \hat{y}^{(m)} - y^{(m)} \end{bmatrix} = \begin{bmatrix} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_1^{(i)} \\ \vdots \\ \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_n^{(i)} \\ \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) \end{bmatrix}$$

(Mini-)Batch Training

$$\begin{bmatrix} x_1^{(1)} & \dots & x_1^{(m)} \\ \vdots & \ddots & \vdots \\ x_n^{(1)} & \dots & x_n^{(m)} \\ 1 & 1 & 1 \end{bmatrix} \cdot \begin{bmatrix} \hat{y}^{(1)} - y^{(1)} \\ \vdots \\ \hat{y}^{(m)} - y^{(m)} \end{bmatrix} = \begin{bmatrix} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_1^{(i)} \\ \vdots \\ \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_n^{(i)} \\ \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) \end{bmatrix}$$
$$= \begin{bmatrix} \sum_{i=1}^m \left(\frac{\partial L}{\partial w_1} \right)^{(i)} \\ \vdots \\ \sum_{i=1}^m \left(\frac{\partial L}{\partial w_n} \right)^{(i)} \\ \sum_{i=1}^m \left(\frac{\partial L}{\partial b} \right)^{(i)} \end{bmatrix}$$

(Mini-)Batch Training

$$\begin{aligned} \begin{bmatrix} x_1^{(1)} & \dots & x_1^{(m)} \\ \vdots & \ddots & \vdots \\ x_n^{(1)} & \dots & x_n^{(m)} \\ 1 & 1 & 1 \end{bmatrix} \cdot \begin{bmatrix} \hat{y}^{(1)} - y^{(1)} \\ \vdots \\ \hat{y}^{(m)} - y^{(m)} \end{bmatrix} &= \begin{bmatrix} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_1^{(i)} \\ \vdots \\ \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_n^{(i)} \\ \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) \end{bmatrix} \\ &= \begin{bmatrix} \sum_{i=1}^m \left(\frac{\partial L}{\partial w_1} \right)^{(i)} \\ \vdots \\ \sum_{i=1}^m \left(\frac{\partial L}{\partial w_n} \right)^{(i)} \\ \sum_{i=1}^m \left(\frac{\partial L}{\partial b} \right)^{(i)} \end{bmatrix} \\ &= \sum_{i=1}^m (\nabla L)^{(i)} \end{aligned}$$

(Mini-)Batch Training

- ▶ $\nabla L = \frac{1}{m} \left(\mathbf{x}^T \cdot (\hat{\mathbf{y}} - \mathbf{y}) \right)$
 - ▶ What is $\mathbf{x}^T \cdot (\hat{\mathbf{y}} - \mathbf{y})$?
 - ▶ It computes the sum of the gradients for each document i in the mini-batch!

(Mini-)Batch Training

- ▶ $\nabla L = \frac{1}{m} \left(\mathbf{x}^T \cdot (\hat{\mathbf{y}} - \mathbf{y}) \right)$
 - ▶ What is $\mathbf{x}^T \cdot (\hat{\mathbf{y}} - \mathbf{y})$?
 - ▶ It computes the sum of the gradients for each document i in the mini-batch!
 - ▶ Then to get the average gradient, we just divide by m